

# Style

November 29, 2019

## 1 Style Transfer

```
[0]: # import resources
    %matplotlib inline

    from PIL import Image
    from io import BytesIO
    import matplotlib.pyplot as plt
    import numpy as np

    import torch
    import torch.optim as optim
    import requests
    from torchvision import transforms, models
```

## 2 VGG19

```
[0]: vgg = models.vgg19(pretrained=True).features

    for param in vgg.parameters():
        param.requires_grad_(False)

[38]: device = torch.device("cuda" if torch.cuda.is_available() else "cpu")

    vgg.to(device)
    print(vgg)
```

```
Sequential(
  (0): Conv2d(3, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (1): ReLU(inplace=True)
  (2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (3): ReLU(inplace=True)
  (4): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1,
    ceil_mode=False)
  (5): Conv2d(64, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
```

```

(6): ReLU(inplace=True)
(7): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(8): ReLU(inplace=True)
(9): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1,
ceil_mode=False)
(10): Conv2d(128, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(11): ReLU(inplace=True)
(12): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(13): ReLU(inplace=True)
(14): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(15): ReLU(inplace=True)
(16): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(17): ReLU(inplace=True)
(18): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1,
ceil_mode=False)
(19): Conv2d(256, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(20): ReLU(inplace=True)
(21): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(22): ReLU(inplace=True)
(23): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(24): ReLU(inplace=True)
(25): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(26): ReLU(inplace=True)
(27): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1,
ceil_mode=False)
(28): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(29): ReLU(inplace=True)
(30): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(31): ReLU(inplace=True)
(32): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(33): ReLU(inplace=True)
(34): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(35): ReLU(inplace=True)
(36): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1,
ceil_mode=False)
)

```

### 3 Load Images

```

[0]: def load_image(img_path, max_size=400, shape=None):
      if "http" in img_path:
          response = requests.get(img_path)
          image = Image.open(BytesIO(response.content)).convert('RGB')
      else:
          image = Image.open(img_path).convert('RGB')

```

```

# large images will slow down processing
if max(image.size) > max_size:
    size = max_size
else:
    size = max(image.size)

if shape is not None:
    size = shape

in_transform = transforms.Compose([
    transforms.Resize(size),
    transforms.ToTensor(),
    transforms.Normalize((0.485, 0.456, 0.406),
                          (0.229, 0.224, 0.225))])

# discard the transparent, alpha channel and add batch
image = in_transform(image)[:3,:,:].unsqueeze(0)

return image

```

```

[0]: def im_convert(tensor):
    image = tensor.to("cpu").clone().detach()
    image = image.numpy().squeeze()
    image = image.transpose(1,2,0)
    image = image * np.array((0.229, 0.224, 0.225)) + np.array((0.485, 0.456, 0.
→406))
    image = image.clip(0, 1)

    return image

```

```

[0]: content = load_image('octopus.jpg').to(device)

style = load_image('Dark.png', shape=content.shape[-2:]).to(device)

```

```

[42]: fig, (ax1, ax2) = plt.subplots(1, 2, figsize=(20, 10))

ax1.imshow(im_convert(content))
ax2.imshow(im_convert(style))

```

```

[42]: <matplotlib.image.AxesImage at 0x7f456e20d4a8>

```



## 4 Content and Style Features

```
[0]: def get_features(image, model, layers=None) -> dict:
      """ Run an image forward through a model and get the features for a set of
      layers. """
      if layers is None:
          layers = {'0': 'conv1_1',
                    '5': 'conv2_1',
                    '10': 'conv3_1',
                    '19': 'conv4_1',
                    '21': 'conv4_2', ## content representation
                    '28': 'conv5_1'}

      features = {}
      x = image

      for name, layer in model._modules.items():
          x = layer(x)
          if name in layers:
              features[layers[name]] = x

      return features
```

## 5 Gram Matrix

```
[0]: def gram_matrix(tensor) -> torch.Tensor:
      """ Calculate the Gram Matrix of a given tensor
      Gram Matrix from :v https://en.wikipedia.org/wiki/Gramian\_matrix
      """

      _, d, h, w = tensor.size()
```

```

# reshape so we're multiplying the features for each channel
tensor = tensor.view(d, h * w)

gram = torch.mm(tensor, tensor.t())

return gram

```

```

[0]: content_features = get_features(content, vgg)
style_features = get_features(style, vgg)

# calculate the gram matrices for each layer of our style representation
style_grams = {layer: gram_matrix(style_features[layer]) for layer in
    style_features}

# create a target image and prep it for change
target = content.clone().requires_grad_(True).to(device)

```

## 6 Loss

```

[0]: style_weights = {'conv1_1': 1.,
                     'conv2_1': 0.75,
                     'conv3_1': 0.2,
                     'conv4_1': 0.2,
                     'conv5_1': 0.2}

content_weight = 1 # alpha
style_weight = 1e6 # beta

```

## 7 Train loop (from Udacity course)

```

[47]: show_every = 350

optimizer = optim.Adam([target], lr=0.003)
steps = 2000

for ii in range(1, steps+1):
    target_features = get_features(target, vgg)
    content_loss = torch.mean((target_features['conv4_2'] -
    content_features['conv4_2'])**2)

    # the style loss
    style_loss = 0

    if ii % 50 == 0: print(f"{ii / steps * 100}%")

```

```

for layer in style_weights:
    # get the "target" style representation for the layer
    target_feature = target_features[layer]
    target_gram = gram_matrix(target_feature)
    _, d, h, w = target_feature.shape

    style_gram = style_grams[layer]
    layer_style_loss = style_weights[layer] * torch.mean((target_gram -
→style_gram)**2)

    style_loss += layer_style_loss / (d * h * w)

total_loss = content_weight * content_loss + style_weight * style_loss

# update target
optimizer.zero_grad()
total_loss.backward()
optimizer.step()


# display intermediate images and print the loss
if ii % show_every == 0:
    print('Total loss: ', total_loss.item())
    plt.imshow(im_convert(target))
    plt.show()

```

```

2.5%
5.0%
7.5%
10.0%
12.5%
15.0%
17.5%
Total loss: 6446763.5

```



20.0%

22.5%

25.0%


27.500000000000004%

30.0%

32.5%

35.0%

Total loss: 1900207.0



37.5%

40.0%

42.5%

45.0%


47.5%

50.0%


52.5%

Total loss: 778403.9375





55.000000000000001%  
57.499999999999999%  
60.0%  
62.5%  
65.0%  
67.5%  
70.0%  
Total loss: 409117.40625



72.5%

75.0%

77.5%


80.0%

82.5%

85.0%

87.5%

Total loss: 256645.5



90.0%  
92.5%  
95.0%  
97.5%  
100.0%

```
[52]: fig, ax1 = plt.subplots(1, 1, figsize=(20, 10))  
      ax1.imshow(im_convert(style))
```

```
[52]: <matplotlib.image.AxesImage at 0x7f456e2f31d0>
```



```
[48]: fig, (ax1, ax2) = plt.subplots(1, 2, figsize=(20, 10))
      ax1.imshow(im_convert(content))
      ax2.imshow(im_convert(target))
```

[48]: <matplotlib.image.AxesImage at 0x7f4566524550>

